

Assignment1

September 8, 2024

1 Programming Assignment 1: MNIST Classification using MLP

1.1 EE5179:Deep Learning for Imaging

1.2 Classification using MLP

Name: Saranath P

Roll No: ns24z458

2 All imports

```
[1]: import numpy as np
from scipy.special import expit
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report
from datetime import datetime
import torch
from torchvision import datasets
from torch.utils.data import DataLoader, random_split, TensorDataset
from torchvision.transforms import ToTensor
```

3 Helper classes and functions.

We will first get some helper classes and functions which will be used for implementing the back-propagation algorithm from scratch

3.1 Loss functions

We define a class on various loss functions

```
[ ]: class LossFunctions:

    class SquaredErrorLoss:
        def __init__(self, alpha=0.01, regularization='l2'):
            """
            Initializes the SquaredErrorLoss with optional L1 or L2
            ↪regularization.
```

```

Parameters:
alpha (float): Regularization constant.
regularization (str): Type of regularization ('l1' for L1, 'l2' for L2, or None).
"""
self.alpha = alpha
self.regularization = regularization

def __call__(self, y_pred, y_true, model_weights=None):
    """
    Computes the Mean Squared Error (MSE) with optional L1 or L2 regularization.

    Parameters:
    y_pred (ndarray): The predicted values.
    y_true (ndarray): The true values.
    model_weights (list): A list of model weights (e.g., numpy arrays).

    Returns:
    float: The regularized mean squared error.
    """
    mse_loss = 0.5 * np.mean((y_pred - y_true) ** 2)

    # Regularization term
    if self.regularization == 'l2':
        reg_term = self.alpha * sum(np.sum(w ** 2) for w in model_weights)
    elif self.regularization == 'l1':
        reg_term = self.alpha * sum(np.sum(np.abs(w)) for w in model_weights)
    else:
        reg_term = 0 # No regularization

    return mse_loss + reg_term

def gradient(self, y_pred, y_true, model_weights=None):
    """
    Computes the gradient of the Mean Squared Error (MSE) Loss with respect to y_pred, including the gradient of the regularization term.

    Parameters:
    y_pred (ndarray): The predicted values.
    y_true (ndarray): The true values.
    model_weights (list): A list of model weights (e.g., numpy arrays).

    Returns:
    """

```

tuple: The gradient with respect to y_pred and the gradients with respect to the model weights.

```
"""
    gradient_y_pred = 2 * (y_pred - y_true)

    gradients_weights = []
    if model_weights is not None:
        for w in model_weights:
            if self.regularization == 'l2':
                grad_w = self.alpha * w
            elif self.regularization == 'l1':
                grad_w = self.alpha * np.sign(w)
            else:
                grad_w = np.zeros_like(w)
            gradients_weights.append(grad_w)

    return gradient_y_pred, gradients_weights

def __repr__(self):
    return f"SquaredErrorLoss(alpha={self.alpha}, regularization={self.
    regularization})"
```

```
class CrossEntropyLoss:
    def __init__(self, alpha=0.01, regularization=None):
        self.alpha = alpha
        self.regularization = regularization

    def __call__(self, y_pred, y_true, model_weights):

        y_pred = np.clip(y_pred, 1e-15, 1 - 1e-15)
        ce_loss = -np.sum(y_true * np.log(y_pred))

        if self.regularization == 'l2':
            reg_term = self.alpha * sum(np.sum(w ** 2) for w in
            model_weights)
        elif self.regularization == 'l1':
            reg_term = self.alpha * sum(np.sum(np.abs(w)) for w in
            model_weights)
        else:
            reg_term = 0

        return ce_loss + reg_term
```

```

def gradient(self, y_pred, y_true, model_weights):
    """
    Compute the gradient of the Cross-Entropy Loss with respect to
    ↪y_pred,
    including the gradient of the regularization term.

    Parameters:
    y_pred (ndarray): The predicted probabilities (output of softmax)
    ↪of shape (num_classes,).
    y_true (ndarray): The one-hot encoded true labels of shape
    ↪(num_classes,).
    model_weights (list): A list of model weights (e.g., numpy arrays).

    Returns:
    list: A list containing the gradient with respect to y_pred and the
    ↪gradients
    with respect to the model weights.
    """

    y_pred = np.clip(y_pred, 1e-15, 1 - 1e-15)
    gradient_y_pred = -(y_true / y_pred)

    gradients_weights = []
    for w in model_weights:
        if self.regularization == 'l2':
            grad_w = self.alpha * w
        elif self.regularization == 'l1':
            grad_w = self.alpha * np.sign(w)
        else:
            grad_w = np.zeros_like(w)
        gradients_weights.append(grad_w)

    return gradient_y_pred, gradients_weights

def __repr__(self):
    return f"CrossEntropyLoss(alpha={self.alpha}, regularization={self.
    ↪regularization})"

```

3.2 Activation functions

We define the activation functions and its gradients.

```
[ ]: class Activations:
```

```

class ReLU:
    def __call__(self, x):
        """
        Applies the ReLU activation function element-wise to the given
        ↪ input x.

        Args:
            x (numpy array): The input array.

        Returns:
            numpy array: The output of the ReLU activation function, which
            ↪ is element-wise  $\max(0, x)$ .
        """
        return np.maximum(0, x)

    def gradient(self, x):
        """
        Computes the gradient (derivative) of the ReLU function for
        ↪ backpropagation.

        Args:
            x (numpy array): The input array.

        Returns:
            numpy array: The gradient of ReLU, which is 1 for  $x > 0$  and 0
            ↪ otherwise.
        """
        return np.where(x > 0, 1, 0)

class Sigmoid:

    def __call__(self, x):
        """
        Applies the Sigmoid activation function to the given input x.

        Args:
            x (float): The input value.

        Returns:
            float: The output of the Sigmoid activation function.
        """
        return expit(x)

    def gradient(self, x):
        """

```

```

        Computes the gradient of the Sigmoid activation function.

    Args:
        x (float): The input value.

    Returns:
        float: The gradient of the Sigmoid activation function.
    """
    return self.__call__(x) * (1 - self.__call__(x))

class Softmax:

    def __call__(self, x):
        """
        Applies the Softmax activation function to the given input x.

    Args:
        x (numpy.ndarray): The input array.

    Returns:
        numpy.ndarray: The output of the Softmax activation function.
        """
        exp_x = np.exp(x - np.max(x)) # Subtracting max(x) for numerical_
↪stability
        return exp_x / np.sum(exp_x, axis=0)

    def gradient(self, x):
        """
        Computes the gradient of the Softmax activation function.

    Args:
        x (numpy.ndarray): The input array.

    Returns:
        numpy.ndarray: The Jacobian matrix of the Softmax activation_
↪function.
        """
        s = self.__call__(x)
        jacobian_matrix = np.diagflat(s) - np.outer(s, s)
        return jacobian_matrix

    def __repr__(self):
        return "Softmax"

```

```

class Tanh:

    def __call__(self, x):
        """
        Applies the Tanh activation function to the given input x.

        Args:
            x (numpy.ndarray): The input array.

        Returns:
            numpy.ndarray: The output of the Tanh activation function.
        """
        return np.tanh(x)

    def gradient(self, x):
        """
        Computes the gradient of the Tanh activation function.

        Args:
            x (numpy.ndarray): The input array.

        Returns:
            numpy.ndarray: The derivative of the Tanh activation function.
        """
        return 1 - np.power(self.__call__(x), 2)

class Linear:

    def __call__(self, x):
        """
        Returns the input value.

        Parameters:
            x (Any): The input value.

        Returns:
            Any: The input value.
        """
        return x

    def gradient(self, x):
        """
        Computes the gradient of the Linear activation function.

        Args:
            x (Any): The input value.

```

```

        Returns:
            int: The derivative of the Linear activation function, which is
↳ always 1.
        """
        return 1

```

3.3 Initialization

We define various initialization methods here

```

[ ]: import numpy as np

class ParameterInitializer:
    def __init__(self, initialization='random'):
        """
        Parameters:
            initialization (str): 'random' for uniform random initialization,
                                'gaussian' for Gaussian distribution,
                                'glorot' for Glorot initialization.
        """
        self.initialization = initialization

    def initialize_parameters(self, inputs, hidden_layers, outputs):
        """
        Initializes the parameters for a neural network.

        Args:
            inputs (int): The number of input nodes.
            hidden_layers (list): A list of integers representing the number of
↳ nodes in each hidden layer.
            outputs (int): The number of output nodes.

        Returns:
            dict: A dictionary containing the initialized parameters, with keys
↳ "W1" and "b1" for the first layer, and "W{i+1}" and "b{i+1}" for each hidden
↳ layer, where i is the index of the layer. The last two keys are
↳ "W{len(hidden_layers) + 1}" and "b{len(hidden_layers) + 1}" for the output
↳ layer.

        Raises:
            None.
        """
        parameters = {}

        if self.initialization == 'random':
            parameters["W1"] = np.random.rand(hidden_layers[0], inputs)

```



```

        parameters["b1"] = np.random.rand(hidden_layers[0])

    elif self.initialization == 'gaussian':
        parameters["W1"] = np.random.randn(hidden_layers[0], inputs)
        parameters["b1"] = np.random.randn(hidden_layers[0])

    elif self.initialization == 'glorot':
        parameters["W1"] = np.random.uniform(-np.sqrt(6 / (inputs +
↪hidden_layers[0])),
                                                np.sqrt(6 / (inputs +
↪hidden_layers[0])),
                                                (hidden_layers[0], inputs))
        parameters["b1"] = np.zeros(hidden_layers[0])

    for i in range(1, len(hidden_layers)):
        if self.initialization == 'random':
            parameters[f"W{i+1}"] = np.random.rand(hidden_layers[i],
↪hidden_layers[i - 1])
            parameters[f"b{i+1}"] = np.random.rand(hidden_layers[i])
        elif self.initialization == 'gaussian':
            parameters[f"W{i+1}"] = np.random.randn(hidden_layers[i],
↪hidden_layers[i - 1])
            parameters[f"b{i+1}"] = np.random.randn(hidden_layers[i])
        elif self.initialization == 'glorot':
            parameters[f"W{i+1}"] = np.random.uniform(-np.sqrt(6 /
↪(hidden_layers[i - 1] + hidden_layers[i])),
                                                        np.sqrt(6 /
↪(hidden_layers[i - 1] + hidden_layers[i])),
                                                        (hidden_layers[i],
↪hidden_layers[i - 1]))
            parameters[f"b{i+1}"] = np.zeros(hidden_layers[i])

    if self.initialization == 'random':
        parameters[f"W{len(hidden_layers) + 1}"] = np.random.rand(outputs,
↪hidden_layers[-1])
        parameters[f"b{len(hidden_layers) + 1}"] = np.random.rand(outputs)
    elif self.initialization == 'gaussian':
        parameters[f"W{len(hidden_layers) + 1}"] = np.random.randn(outputs,
↪hidden_layers[-1])
        parameters[f"b{len(hidden_layers) + 1}"] = np.random.randn(outputs)
    elif self.initialization == 'glorot':
        parameters[f"W{len(hidden_layers) + 1}"] = np.random.uniform(-np.
↪sqrt(6 / (hidden_layers[-1] + outputs)),
                                                                    np.
↪sqrt(6 / (hidden_layers[-1] + outputs)),

```

```

    ↪(outputs, hidden_layers[-1]))
        parameters[f"b{len(hidden_layers) + 1}"] = np.zeros(outputs)

    return parameters

```

4 Feed Forward Neural Netowrk Implementation

```

[ ]: softmax = Activations.Softmax()
    sigmoid = Activations.Sigmoid()
    crossentropy = LossFunctions.CrossEntropyLoss()

class FeedForwardNeuralNets:
    def __init__(self, inputs, hidden_layers, outputs, g=sigmoid, ↪
    ↪L=crossentropy, O=softmax, eta=0.01,
        optimizer="gd", initialization_method="glorot", batch_size=32,
        beta1=0.9, beta2=0.999, epsilon=1e-8, t=0):
        self.inputs = inputs
        self.outputs = outputs
        self.batch_size = min(batch_size, inputs.shape[0])
        if len(self.outputs.shape) < 2:
            self.parameters = ParameterInitializer(initialization_method).
            ↪initialize_parameters(
                inputs[0].shape[0], hidden_layers, 1)
        else:
            self.parameters = ParameterInitializer(initialization_method).
            ↪initialize_parameters(
                inputs[0].shape[0], hidden_layers, outputs[0].shape[0])
        self.g = g
        self.O = O
        self.L = L
        self.eta = eta
        self.losses = {}
        self.activations = {}
        self.beta1 = beta1
        self.beta2 = beta2
        self.epsilon = epsilon
        self.t = t
        self.v = {key: np.zeros_like(value) for key, value in self.parameters.
            ↪items()}
        self.s = {key: np.zeros_like(value) for key, value in self.parameters.
            ↪items()}
        self.training_loss = []
        self.validation_loss = []

```

```

if optimizer == "gd":
    self.optimizer = self.gradient_descent
elif optimizer == "adam":
    self.optimizer = self.adam

def forward_propagation(self, x):
    """
    Performs forward propagation through the neural network.

    Parameters:
        x (numpy array): The input to the neural network.

    Returns:
        y_pred (numpy array): The predicted output of the neural network.
    """

    self.activations["a1"] = self.parameters["W1"] @ x + self.
    ↪parameters["b1"]
    self.activations["h1"] = self.g(self.activations["a1"])
    for i in range(2, len(self.parameters) // 2):
        self.activations[f"a{i}"] = self.parameters[f"W{i}"] @ self.
    ↪activations[f"h{i - 1}"] + self.parameters[f"b{i}"]
        self.activations[f"h{i}"] = self.g(self.activations[f"a{i}"])

    self.activations[f"a{len(self.parameters) // 2}"] = self.
    ↪parameters[f"W{len(self.parameters) // 2}"] @ self.activations[f"h{len(self.
    ↪parameters) // 2 - 1}"] + self.parameters[f"b{len(self.parameters) // 2}"]
    y_pred = self.0(self.activations[f"a{len(self.parameters) // 2}"])
    return y_pred

def backPropogation(self, y_pred, y, x):
    """
    Performs backpropagation through the neural network.

    Parameters:
        y_pred (numpy array): The predicted output of the neural network.
        y (numpy array): The actual output of the neural network.
        x (numpy array): The input to the neural network.

    Returns:
        None
    """

    n = len(self.parameters) // 2
    m = len(self.activations) // 2

    La = y_pred - y

```

```

Lh = La @ self.parameters[f"W{n}"]

da = self.g.gradient(self.activations[f"a{m}"])
self.losses[f"W{n}"] = np.outer(La, self.activations[f"h{m}"])
self.losses[f"b{n}"] = La.copy()

for i in range(1, m):
    La = Lh * da
    Lh = La @ self.parameters[f"W{m - i + 1}"]
    da = self.g.gradient(self.activations[f"a{m - i}"])
    self.losses[f"W{m - i + 1}"] = np.outer(La, self.activations[f"h{m-
↪- i}"])
    self.losses[f"b{m - i + 1}"] = La.copy()

La = Lh * da
self.losses["W1"] = np.outer(La, x)
self.losses["b1"] = La.copy()

_, gradients_weights = self.L.gradient(y_pred, y, list(self.parameters.
↪values()))

for i, key in enumerate(self.parameters.keys()):
    if key.startswith('W'):
        self.losses[key] += gradients_weights[i]

def gradient_descent(self):
    """
    Performs gradient descent optimization on the neural network's
    ↪parameters.

    Updates the weights (W) and biases (b) of the neural network by
    ↪subtracting the product of the learning rate (eta) and the corresponding
    ↪losses.

    Parameters:
        None

    Returns:
        None
    """
    for i in range(len(self.parameters) // 2):
        self.parameters[f"W{i + 1}"] = self.parameters[f"W{i + 1}"] - self.
↪eta * self.losses[f"W{i + 1}"]
        self.parameters[f"b{i + 1}"] = self.parameters[f"b{i + 1}"] - self.
↪eta * self.losses[f"b{i + 1}"]

```

```

def adam(self):
    """
        Performs Adam optimization on the neural network's parameters.

        Updates the weights (W) and biases (b) of the neural network using the
        Adam optimization algorithm.

        Parameters:
            None

        Returns:
            None
    """
    self.t += 1
    for key in self.parameters.keys():
        if key.startswith("W") or key.startswith("b"):
            gradient = self.losses[key]

            self.v[key] = self.beta1 * self.v[key] + (1 - self.beta1) *
            gradient
            self.s[key] = self.beta2 * self.s[key] + (1 - self.beta2) *
            (gradient ** 2)

            v_corrected = self.v[key] / (1 - self.beta1 ** self.t)
            s_corrected = self.s[key] / (1 - self.beta2 ** self.t)

            self.parameters[key] -= self.eta * v_corrected / (np.
            sqrt(s_corrected) + self.epsilon)

def train(self, epochs, validation_data=None):
    """
        Trains the neural network model for a specified number of epochs.

        This function takes in the number of epochs as input and performs the
        following operations:
        - Shuffles the input data and corresponding outputs.
        - Divides the shuffled data into batches based on the batch size.
        - For each batch, it calculates the total gradient loss by
        iterating over each input and output pair.
        - It then updates the model parameters using the optimizer function.
        - If validation data is provided, calculates and appends the
        validation loss every 200 iterations.

        Parameters:
    """

```

epochs (int): The number of epochs to train the model for.
validation_data (tuple): A tuple containing validation inputs and
→ outputs.

Returns:

None

"""

```

for epoch in range(epochs):
    permutation = np.random.permutation(self.inputs.shape[0])
    inputs_shuffled = self.inputs[permutation]
    outputs_shuffled = self.outputs[permutation]

    total_loss = 0

    for i in range(0, self.inputs.shape[0], self.batch_size):
        batch_inputs = inputs_shuffled[i:i + self.batch_size]
        batch_outputs = outputs_shuffled[i:i + self.batch_size]

        total_gradient_loss = {key: 0 for key in self.parameters.keys()}
        batch_loss = 0

        for x, y in zip(batch_inputs, batch_outputs):
            y_pred = self.forward_propagation(x)
            batch_loss += self.L(y_pred, y, list(self.parameters.
→ values()))

            self.backPropogation(y_pred, y, x)

            for key in total_gradient_loss:
                total_gradient_loss[key] += self.losses.get(key, 0)

        for key in total_gradient_loss:
            total_gradient_loss[key] /= self.batch_size

        self.optimizer()

        total_loss += batch_loss / self.batch_size

    if (i // self.batch_size) % 200 == 0:
        iteration_loss = batch_loss / self.batch_size
        self.training_loss.append(iteration_loss)

    if validation_data is not None:
        val_inputs, val_outputs = validation_data
        val_loss = 0
        for x_val, y_val in zip(val_inputs, val_outputs):
            y_pred_val = self.forward_propagation(x_val)

```

```

        val_loss += self.L(y_pred_val, y_val, list(self.
↪parameters.values()))
        val_loss /= val_inputs.shape[0]
        self.validation_loss.append(val_loss)

    epoch_loss = total_loss / (self.inputs.shape[0] // self.batch_size)
    if validation_data is not None:
        val_loss = self.validation_loss[-1] if self.validation_loss_
↪else None
        print(f'Epoch {epoch + 1}/{epochs}, Training Loss: {epoch_loss:.
↪4f}, Validation Loss: {val_loss:.4f}')
    else:
        print(f'Epoch {epoch + 1}/{epochs}, Training Loss: {epoch_loss:.
↪4f}')

def evaluate(self, X):
    """
    Evaluates the neural network model for the given input X.

    Parameters:
        X (numpy array): The input to the neural network.

    Returns:
        numpy array: The predicted output of the neural network.
    """
    return np.array(self.forward_propagation(X))

def predict(self, X):
    """
    Uses evaluate function and get arg max and get the prediction

    Parameters:
        X (numpy array): The input to the neural network

    Returns:
        Class label: The label that has the highest probabilitiy
    """
    return np.argmax(self.evaluate(X))

```

5 Example Implementation

```

[ ]: X = np.array([
    [0.1, 0.5, 0.2, 0.9, 0.3, 0.6, 0.4, 0.8, 0.7, 0.2],
    [0.3, 0.2, 0.7, 0.4, 0.6, 0.8, 0.9, 0.5, 0.1, 0.3],
    [0.8, 0.9, 0.4, 0.1, 0.5, 0.2, 0.7, 0.6, 0.3, 0.9],

```

```

    [0.6, 0.2, 0.3, 0.7, 0.4, 0.1, 0.8, 0.9, 0.5, 0.6],
    [0.5, 0.1, 0.8, 0.2, 0.9, 0.7, 0.3, 0.4, 0.6, 0.5],
    [0.7, 0.3, 0.6, 0.5, 0.8, 0.4, 0.1, 0.2, 0.9, 0.7],
    [0.2, 0.8, 0.1, 0.6, 0.4, 0.5, 0.7, 0.3, 0.2, 0.8],
    [0.9, 0.6, 0.5, 0.3, 0.7, 0.2, 0.4, 0.1, 0.8, 0.6],
    [0.4, 0.7, 0.9, 0.8, 0.2, 0.3, 0.6, 0.5, 0.1, 0.4],
    [0.5, 0.3, 0.7, 0.4, 0.1, 0.8, 0.2, 0.6, 0.7, 0.5]
])

y = np.array([1, 0, 1, 0, 1, 0, 1, 0, 1, 0])
num_classes = len(np.unique(y))
y_one_hot = np.eye(num_classes)[y]

print("One-hot encoded y:")
print(y_one_hot)

```

One-hot encoded y:

```

[[0. 1.]
 [1. 0.]
 [0. 1.]
 [1. 0.]
 [0. 1.]
 [1. 0.]
 [0. 1.]
 [1. 0.]
 [0. 1.]
 [1. 0.]]

```

```

[ ]: g_1 = LossFunctions.CrossEntropyLoss()
      g_2 = LossFunctions.CrossEntropyLoss(regularization='l2')

      model = FeedForwardNeuralNets(X, [50], y_one_hot, L=g_2, optimizer="adam",
      ↪g=Activations.ReLU(), eta=0.01,
      )
      model.train(epochs=500, validation_data=[X, y_one_hot])

```

```

[ ]: for i in X:
      print(model.evaluate(i))

```

```

[0.07989294 0.92010706]
[0.77283669 0.22716331]
[0.04117688 0.95882312]
[0.93341553 0.06658447]
[0.34602394 0.65397606]
[0.94823665 0.05176335]
[0.0046353 0.9953647]
[0.93805742 0.06194258]
[0.04854306 0.95145694]

```



```
[0.76807655 0.23192345]
```

```
[ ]: model.parameters
```


6 Importing MNIST Dataset

```
[2]: training_data = datasets.MNIST(  
    root="data",  
    train=True,  
    download=True,  
    transform=ToTensor()  
)  
  
test_data = datasets.MNIST(  
    root="data",  
    train=False,  
    download=True,  
    transform=ToTensor()  
)
```

```
[3]: validation_size = 10000  
training_size = len(training_data) - validation_size  
  
train_dataset, val_dataset = random_split(training_data, [training_size, validation_size])  
  
train_loader = DataLoader(training_data, batch_size=training_size, shuffle=True)  
val_loader = DataLoader(val_dataset, batch_size=validation_size, shuffle=False)  
  
train_data = next(iter(train_loader))  
X_train = train_data[0].numpy()  
y_train = train_data[1].numpy()  
  
val_data = next(iter(val_loader))  
X_val = val_data[0].numpy()  
y_val = val_data[1].numpy()
```

```
[ ]: plt.imshow(X_train[6].reshape(28, 28), cmap="gray"), y_train[6]
```

```
[ ]: (<matplotlib.image.AxesImage at 0x7f9c7aaad280>, 0)
```



7 Normalize the inputs to zero mean and unit variance

```
[ ]: def preprocess(X, y):  
  
    X = X / 255.0  
    mean = X.mean()  
    std = X.std()  
    X = (X - mean) / std  
  
    X = X.reshape(X.shape[0], -1)  
  
    num_classes = len(np.unique(y))  
    return X, np.eye(num_classes)[y]  
  
X_train, y_train_one_hot = preprocess(X_train, y_train)  
X_val, y_val_one_hot = preprocess(X_val, y_val)
```

8 Train the network with given configuration

As given as the base line model we will train the network according to the given configuration.

Default : - $\eta = 0.01$ - Initialization method : Glorot - Activation : Sigmoid - Loss function : Cross Entropy - Output function : Softmax

Hidden Layers : - $h_1 = 500$ - $h_2 = 250$ - $h_3 = 100$ - $i/p = 28 \times 28 = 764 \times 1$ (After flattening) - $o/p = 10 \times 1$ (Number of classes)


Number of epochs : 15

```
[ ]: baselineModel = FeedForwardNeuralNets(X_train, [500, 250, 100],  
      ↪y_train_one_hot, optimizer="gd", batch_size=64)  
  
start_time = datetime.now()  
  
baselineModel.train(epochs=15, validation_data=[X_val, y_val_one_hot])  
  
end_time = datetime.now()  
  
time_taken = end_time - start_time  
print(f"Total time taken for training: {time_taken}")  
  
# This cell takes to run approx. 44 mins with a 8 core CPU
```

```
Epoch 1/15, Training Loss: 2.3346, Validation Loss: 2.3094  
Epoch 2/15, Training Loss: 2.2171, Validation Loss: 2.1922  
Epoch 3/15, Training Loss: 1.9972, Validation Loss: 1.8870  
Epoch 4/15, Training Loss: 1.5686, Validation Loss: 1.4182  
Epoch 5/15, Training Loss: 1.1994, Validation Loss: 1.1047  
Epoch 6/15, Training Loss: 0.9789, Validation Loss: 0.9207  
Epoch 7/15, Training Loss: 0.8371, Validation Loss: 0.7936  
Epoch 8/15, Training Loss: 0.7298, Validation Loss: 0.7354  
Epoch 9/15, Training Loss: 0.6592, Validation Loss: 0.6320  
Epoch 10/15, Training Loss: 0.5953, Validation Loss: 0.5839  
Epoch 11/15, Training Loss: 0.5517, Validation Loss: 0.5399  
Epoch 12/15, Training Loss: 0.5212, Validation Loss: 0.5037  
Epoch 13/15, Training Loss: 0.4858, Validation Loss: 0.5226  
Epoch 14/15, Training Loss: 0.4673, Validation Loss: 0.4718  
Epoch 15/15, Training Loss: 0.4487, Validation Loss: 0.4499  
Total time taken for training: 0:44:24.319133
```

```
[ ]: plt.imshow(X_val[6].reshape(28, 28))
```

```
[ ]: <matplotlib.image.AxesImage at 0x7f9c78739490>
```



```
[ ]: baselineModel.predict(X_val[600]), y_val[600]
```

```
[ ]: (5, 5)
```

9 Plotting the Train and validation loss for every 200 iterations for baseline model


We will try to plot how the loss varies for every 200 iterations.

```
[ ]: x_axis = np.arange(60)
train_loss = np.array(baselineModel.training_loss)
validation_loss = np.array(baselineModel.validation_loss)

plt.plot(x_axis, train_loss, label='Training Loss')
plt.plot(x_axis, validation_loss, label='Validation Loss')

plt.title('Training vs Validation Loss')
plt.xlabel('Iteration (in 200s)')
plt.ylabel('Loss')
```

```
plt.legend()
plt.savefig('training_vs_validation_loss.png')
plt.show()
```



10 Testing with actual test data

```
[ ]: test_loader = DataLoader(test_data, batch_size=len(test_data), shuffle=False)
```

```
X_test, y_test = next(iter(test_loader))
X_test = X_test.numpy()
y_test = y_test.numpy()

X_test, y_test_one_hot = preprocess(X_test, y_test)
```

```
[ ]: X_train.shape, X_test.shape
```

```
[ ]: ((50000, 784), (10000, 784))
```

```
[ ]: baselineModel.predict(X_test[6]), y_test[6]
```

```
[ ]: (4, 4)
```

11 Calculate Loss for test data and plot it


```
[ ]: test_losses = []

for x, y_true in zip(X_test, y_test_one_hot):
    y_pred = baselineModel.forward_propagation(x)
    loss = baselineModel.L(y_pred, y_true, list(baselineModel.parameters))
    test_losses.append(loss)

test_losses = np.array(test_losses)
test_losses
```

```
[ ]: array([0.01107608, 0.07231415, 0.03132954, ..., 0.90026971, 0.14412548,
          0.10986916])
```

```
[ ]: plt.plot(test_losses)
plt.title('Test Loss for Each Sample')
plt.xlabel('Sample Index')
plt.ylabel('Loss')
plt.show()
```



12 Classification Report

12.1 Train Data

```
[ ]: y_train_pred = []
     for x in X_train:
         y_pred = baselineModel.predict(x)
         y_train_pred.append(y_pred)

     y_train_pred = np.array(y_train_pred)
     print(classification_report(y_train, y_train_pred))
```

	precision	recall	f1-score	support
0	0.90	0.97	0.93	4968
1	0.87	0.98	0.92	5689
2	0.86	0.86	0.86	4967
3	0.84	0.86	0.85	5141
4	0.92	0.78	0.84	4872
5	0.85	0.79	0.82	4524

6	0.97	0.86	0.91	4932
7	0.72	0.96	0.83	5192
8	0.90	0.77	0.83	4832
9	0.77	0.68	0.72	4883
accuracy			0.85	50000
macro avg	0.86	0.85	0.85	50000
weighted avg	0.86	0.85	0.85	50000

12.2 Test data

```
[ ]: y_test_pred = []
      for x in X_test:
          y_pred = baselineModel.predict(x)
          y_test_pred.append(y_pred)

      y_test_pred = np.array(y_test_pred)
      print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
0	0.89	0.98	0.93	980
1	0.90	0.99	0.94	1135
2	0.87	0.86	0.86	1032
3	0.83	0.88	0.86	1010
4	0.90	0.78	0.84	982
5	0.85	0.77	0.81	892
6	0.97	0.84	0.90	958
7	0.73	0.95	0.83	1028
8	0.91	0.77	0.84	974
9	0.79	0.72	0.76	1009
accuracy			0.86	10000
macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000

13 Using different activation functions to train the model

We will use two activation functions mentioned below and train our model

- Tanh
- ReLU

13.1 Using Tanh

```
[ ]: model_tanh = FeedForwardNeuralNets(X_train, [500, 250, 100], y_train_one_hot,
    ↪optimizer="gd", batch_size=64, g=Activations.Tanh())

start_time = datetime.now()

model_tanh.train(epochs=15, validation_data=[X_val, y_val_one_hot])

end_time = datetime.now()

time_taken = end_time - start_time
print(f"Total time taken for training: {time_taken}")

# This cell runs for 57 mins
```

```
Epoch 1/15, Training Loss: 0.8020, Validation Loss: 0.6118
Epoch 2/15, Training Loss: 0.4991, Validation Loss: 0.6735
Epoch 3/15, Training Loss: 0.4243, Validation Loss: 0.3930
Epoch 4/15, Training Loss: 0.3991, Validation Loss: 0.4073
Epoch 5/15, Training Loss: 0.3590, Validation Loss: 0.3651
Epoch 6/15, Training Loss: 0.3477, Validation Loss: 0.3610
Epoch 7/15, Training Loss: 0.3028, Validation Loss: 0.2737
Epoch 8/15, Training Loss: 0.2937, Validation Loss: 0.2821
Epoch 9/15, Training Loss: 0.2928, Validation Loss: 0.2897
Epoch 10/15, Training Loss: 0.2627, Validation Loss: 0.2608
Epoch 11/15, Training Loss: 0.2649, Validation Loss: 0.2623
Epoch 12/15, Training Loss: 0.2462, Validation Loss: 0.2625
Epoch 13/15, Training Loss: 0.2376, Validation Loss: 0.2231
Epoch 14/15, Training Loss: 0.2520, Validation Loss: 0.2511
Epoch 15/15, Training Loss: 0.2275, Validation Loss: 0.2076
Total time taken for training: 0:57:37.176801
```

13.2 Using ReLU

```
[ ]: model_relu = FeedForwardNeuralNets(X_train, [500, 250, 100], y_train_one_hot,
    ↪optimizer="gd", batch_size=64, g=Activations.ReLU())

start_time = datetime.now()

model_relu.train(epochs=15, validation_data=[X_val, y_val_one_hot])
```

```

end_time = datetime.now()

time_taken = end_time - start_time
print(f"Total time taken for training: {time_taken}")

# This cell takes 57 mins 53 seconds in total

```

```

Epoch 1/15, Training Loss: 0.9837, Validation Loss: 0.6321
Epoch 2/15, Training Loss: 0.5767, Validation Loss: 0.4795
Epoch 3/15, Training Loss: 0.4467, Validation Loss: 0.4052
Epoch 4/15, Training Loss: 0.4186, Validation Loss: 0.3829
Epoch 5/15, Training Loss: 0.4126, Validation Loss: 0.4309
Epoch 6/15, Training Loss: 0.3590, Validation Loss: 0.3217
Epoch 7/15, Training Loss: 0.3294, Validation Loss: 0.3133
Epoch 8/15, Training Loss: 0.3318, Validation Loss: 0.3401
Epoch 9/15, Training Loss: 0.2767, Validation Loss: 0.2654
Epoch 10/15, Training Loss: 0.3018, Validation Loss: 0.2672
Epoch 11/15, Training Loss: 0.2871, Validation Loss: 0.2572
Epoch 12/15, Training Loss: 0.2847, Validation Loss: 0.3320
Epoch 13/15, Training Loss: 0.2931, Validation Loss: 0.2698
Epoch 14/15, Training Loss: 0.2455, Validation Loss: 0.2559
Epoch 15/15, Training Loss: 0.2233, Validation Loss: 0.2131
Total time taken for training: 0:57:53.821745

```

14 Plotting the Train and validation loss for every 200 iterations for Tanh model

We will try to plot how the loss varies for every 200 iterations.

```


[ ]: x_axis = np.arange(60)
train_loss = np.array(model_tanh.training_loss)
validation_loss = np.array(model_tanh.validation_loss)

plt.plot(x_axis, train_loss, label='Training Loss')
plt.plot(x_axis, validation_loss, label='Validation Loss')

plt.title('Training vs Validation Loss for Tanh model')
plt.xlabel('Iteration (in 200s)')
plt.ylabel('Loss')

plt.legend()
plt.savefig('training_vs_validation_loss_tanh.png')
plt.show()

```



15 Plotting the Train and validation loss for every 200 iterations for ReLU model

We will try to plot how the loss varies for every 200 iterations.


```
[ ]: x_axis = np.arange(60)
train_loss = np.array(model_relu.training_loss)
validation_loss = np.array(model_relu.validation_loss)

plt.plot(x_axis, train_loss, label='Training Loss')
plt.plot(x_axis, validation_loss, label='Validation Loss')

plt.title('Training vs Validation Loss for ReLU model')
plt.xlabel('Iteration (in 200s)')
plt.ylabel('Loss')

plt.legend()
plt.savefig('training_vs_validation_loss_relu.png')
```

```
plt.show()
```



16 Classification report for Tanh and ReLU

16.1 Tanh

Train Data

```
[ ]: y_train_pred_tanh = []
for x in X_train:
    y_pred = model_tanh.predict(x)
    y_train_pred_tanh.append(y_pred)

y_train_pred_tanh = np.array(y_train_pred_tanh)
print(classification_report(y_train, y_train_pred_tanh))
```

	precision	recall	f1-score	support
0	0.96	0.97	0.96	4968
1	0.98	0.98	0.98	5689

2	0.93	0.95	0.94	4967
3	0.94	0.90	0.92	5141
4	0.96	0.89	0.92	4872
5	0.95	0.89	0.92	4524
6	0.98	0.93	0.95	4932
7	0.97	0.90	0.93	5192
8	0.85	0.94	0.89	4832
9	0.82	0.96	0.88	4883
accuracy				0.93 50000
macro avg				0.93 0.93 0.93 50000
weighted avg				0.93 0.93 0.93 50000

16.2 Tanh

Test Data

```
[ ]: y_test_pred_tanh = []
for x in X_test:
    y_pred = model_tanh.predict(x)
    y_test_pred_tanh.append(y_pred)

y_test_pred_tanh = np.array(y_test_pred_tanh)
print(classification_report(y_test, y_test_pred_tanh))
```

	precision	recall	f1-score	support
0	0.94	0.99	0.96	980
1	0.98	0.98	0.98	1135
2	0.92	0.94	0.93	1032
3	0.93	0.91	0.92	1010
4	0.96	0.89	0.93	982
5	0.94	0.89	0.91	892
6	0.98	0.91	0.94	958
7	0.95	0.88	0.92	1028
8	0.87	0.93	0.90	974
9	0.83	0.95	0.89	1009
accuracy				0.93 10000
macro avg				0.93 0.93 0.93 10000
weighted avg				0.93 0.93 0.93 10000

16.3 ReLU

Train data

```
[ ]: y_train_pred_relu = []
      for x in X_train:
          y_pred = model_relu.predict(x)
          y_train_pred_relu.append(y_pred)

      y_train_pred_relu = np.array(y_train_pred_relu)
      print(classification_report(y_train, y_train_pred_relu))
```

	precision	recall	f1-score	support
0	0.93	0.99	0.96	4968
1	0.99	0.97	0.98	5689
2	0.91	0.96	0.94	4967
3	0.96	0.92	0.94	5141
4	0.94	0.95	0.94	4872
5	0.88	0.95	0.91	4524
6	0.96	0.96	0.96	4932
7	0.96	0.93	0.95	5192
8	0.95	0.88	0.91	4832
9	0.92	0.91	0.91	4883
accuracy			0.94	50000
macro avg	0.94	0.94	0.94	50000
weighted avg	0.94	0.94	0.94	50000

```
[ ]: y_test_pred_relu = []
      for x in X_test:
          y_pred = model_relu.predict(x)
          y_test_pred_relu.append(y_pred)

      y_test_pred_relu = np.array(y_test_pred_relu)
      print(classification_report(y_test, y_test_pred_relu))

      # The output is there in the report. Please look at that
```

17 Regularization on the better model

We see that ReLU performs better than sigmoid and Tanh. So we will perform L_2 regularization on model ReLU

```
[ ]: cross_entropy_reg_l2 = LossFunctions.CrossEntropyLoss(regularization='l2',
    ↪alpha=0.01)

      model_relu_l2 = FeedForwardNeuralNets(X_train, [500, 250, 100],
    ↪y_train_one_hot, optimizer="gd", batch_size=64, g=Activations.ReLU(),
    ↪L=cross_entropy_reg_l2)
```

```

start_time = datetime.now()

model_relu_l2.train(epochs=15, validation_data=[X_val, y_val_one_hot])

end_time = datetime.now()

time_taken = end_time - start_time
print(f"Total time taken for training: {time_taken}")

# Total time taken : 1 hr 06 mins

```

```

Epoch 1/15, Training Loss: 11.3984, Validation Loss: 10.7624
Epoch 2/15, Training Loss: 9.5719, Validation Loss: 9.1938
Epoch 3/15, Training Loss: 8.3015, Validation Loss: 8.1260
Epoch 4/15, Training Loss: 7.2534, Validation Loss: 7.4504
Epoch 5/15, Training Loss: 6.3233, Validation Loss: 6.1156
Epoch 6/15, Training Loss: 5.5729, Validation Loss: 5.4898
Epoch 7/15, Training Loss: 4.9239, Validation Loss: 4.7039
Epoch 8/15, Training Loss: 4.3336, Validation Loss: 4.2440
Epoch 9/15, Training Loss: 3.8721, Validation Loss: 3.6975
Epoch 10/15, Training Loss: 3.4169, Validation Loss: 3.2924
Epoch 11/15, Training Loss: 3.0496, Validation Loss: 2.9853
Epoch 12/15, Training Loss: 2.7812, Validation Loss: 2.7342
Epoch 13/15, Training Loss: 2.5181, Validation Loss: 2.4351
Epoch 14/15, Training Loss: 2.2807, Validation Loss: 2.2350
Epoch 15/15, Training Loss: 2.0902, Validation Loss: 2.0112
Total time taken for training: 1:03:26.456915

```

18 Plotting graphs and getting classification report

```

[ ]: x_axis = np.arange(60)
train_loss = np.array(model_relu_l2.training_loss)
validation_loss = np.array(model_relu_l2.validation_loss)


plt.plot(x_axis, train_loss, label='Training Loss')
plt.plot(x_axis, validation_loss, label='Validation Loss')

plt.title('Training vs Validation Loss for ReLU with $L_2$ regularization model,
↪(alpha = 0.01)')
plt.xlabel('Iteration (in 200s)')
plt.ylabel('Loss')

```

```
plt.legend()
plt.savefig('training_vs_validation_loss_relu_l2.png')
plt.show()
```

Training vs Validation Loss for ReLU with L_2 regularization model (alpha = 0.01)



```
[ ]: y_train_pred_relu_l2 = []
for x in X_train:
    y_pred = model_relu_l2.predict(x)
    y_train_pred_relu_l2.append(y_pred)

y_train_pred_relu_l2 = np.array(y_train_pred_relu_l2)
print(classification_report(y_train, y_train_pred_relu_l2))
```

	precision	recall	f1-score	support
0	0.89	0.94	0.92	4968
1	0.65	0.98	0.79	5689
2	0.87	0.72	0.79	4967
3	0.89	0.62	0.73	5141
4	0.77	0.85	0.81	4872
5	0.73	0.78	0.76	4524
6	0.82	0.92	0.87	4932
7	0.72	0.89	0.80	5192
8	0.78	0.58	0.66	4832

	9	0.84	0.47	0.60	4883
accuracy				0.78	50000
macro avg		0.80	0.78	0.77	50000
weighted avg		0.79	0.78	0.77	50000

19 Experiments with different models

```
[ ]: simple_model = FeedForwardNeuralNets(X_train, [60, 60], y_train_one_hot,
    ↪optimizer='adam', batch_size=128, g=Activations.ReLU())

start_time = datetime.now()

simple_model.train(epochs=15, validation_data=[X_val, y_val_one_hot])

end_time = datetime.now()

time_taken = end_time - start_time
print(f"Total time taken for training: {time_taken}")

# Total time taken : 5 min 54 sec
```

```
Epoch 1/15, Training Loss: 2.0271, Validation Loss: 2.0511
Epoch 2/15, Training Loss: 1.6989, Validation Loss: 1.3248
Epoch 3/15, Training Loss: 1.9417, Validation Loss: 2.5326
Epoch 4/15, Training Loss: 1.7098, Validation Loss: 1.4595
Epoch 5/15, Training Loss: 1.5395, Validation Loss: 1.3042
Epoch 6/15, Training Loss: 1.4193, Validation Loss: 1.4811
Epoch 7/15, Training Loss: 1.3879, Validation Loss: 1.3884
Epoch 8/15, Training Loss: 1.2468, Validation Loss: 1.4849
Epoch 9/15, Training Loss: 1.3437, Validation Loss: 1.3594
Epoch 10/15, Training Loss: 1.5493, Validation Loss: 1.4520
Epoch 11/15, Training Loss: 1.6902, Validation Loss: 1.5791
Epoch 12/15, Training Loss: 1.5291, Validation Loss: 1.6458
Epoch 13/15, Training Loss: 1.6245, Validation Loss: 1.8170
Epoch 14/15, Training Loss: 1.6224, Validation Loss: 1.7843
Epoch 15/15, Training Loss: 1.8279, Validation Loss: 1.4440
Total time taken for training: 0:05:54.533827
```

```
[ ]: x_axis = np.arange(30)
train_loss = np.array(simple_model.training_loss)
validation_loss = np.array(simple_model.validation_loss)
```


```

plt.plot(x_axis, train_loss, label='Training Loss')
plt.plot(x_axis, validation_loss, label='Validation Loss')

plt.title('Training vs Validation Loss for 60 x 60 hidden layer model')
plt.xlabel('Iteration (in 200s)')
plt.ylabel('Loss')

plt.legend()
plt.savefig('training_vs_validation_loss_rsimple.png')
plt.show()

```



```

[ ]: y_train_pred_simple = []
for x in X_train:
    y_pred = simple_model.predict(x)
    y_train_pred_simple.append(y_pred)

y_train_pred_simple = np.array(y_train_pred_simple)
print(classification_report(y_train, y_train_pred_simple))

```

	precision	recall	f1-score	support
0	0.98	0.66	0.79	4968
1	0.24	0.99	0.39	5689
2	0.90	0.55	0.68	4967
3	0.87	0.36	0.51	5141
4	0.47	0.79	0.59	4872
5	0.03	0.00	0.00	4524
6	1.00	0.12	0.21	4932
7	0.96	0.72	0.83	5192
8	0.01	0.00	0.00	4832
9	0.20	0.18	0.19	4883
accuracy			0.45	50000
macro avg	0.57	0.44	0.42	50000
weighted avg	0.57	0.45	0.43	50000

```
[ ]: simple_model_tanh = FeedForwardNeuralNets(X_train, [75, 100, 60],
    ↪ y_train_one_hot, optimizer='adam', batch_size=128, g=Activations.Tanh())

start_time = datetime.now()

simple_model_tanh.train(epochs=15, validation_data=[X_val, y_val_one_hot])

end_time = datetime.now()

time_taken = end_time - start_time
print(f"Total time taken for training: {time_taken}")

# Total time taken :
```

```
Epoch 1/15, Training Loss: 2.8796, Validation Loss: 2.7155
Epoch 2/15, Training Loss: 2.7369, Validation Loss: 2.4332
Epoch 3/15, Training Loss: 2.7762, Validation Loss: 3.0455
Epoch 4/15, Training Loss: 2.8017, Validation Loss: 2.5939
Epoch 5/15, Training Loss: 2.6731, Validation Loss: 2.6385
Epoch 6/15, Training Loss: 2.7330, Validation Loss: 2.6124
Epoch 7/15, Training Loss: 2.8385, Validation Loss: 3.1715
Epoch 8/15, Training Loss: 2.8315, Validation Loss: 2.9605
Epoch 9/15, Training Loss: 2.7560, Validation Loss: 2.6486
Epoch 10/15, Training Loss: 2.7437, Validation Loss: 2.6045
Epoch 11/15, Training Loss: 2.7349, Validation Loss: 2.6444
Epoch 12/15, Training Loss: 2.7772, Validation Loss: 2.7866
Epoch 13/15, Training Loss: 2.8654, Validation Loss: 3.2698
```


Epoch 14/15, Training Loss: 2.7987, Validation Loss: 3.0186
Epoch 15/15, Training Loss: 2.7348, Validation Loss: 2.8175
Total time taken for training: 0:07:12.886348

```
[ ]: x_axis = np.arange(30)
train_loss = np.array(simple_model_tanh.training_loss)
validation_loss = np.array(simple_model_tanh.validation_loss)

plt.plot(x_axis, train_loss, label='Training Loss')
plt.plot(x_axis, validation_loss, label='Validation Loss')

plt.title('Training vs Validation Loss for 75 x 100 x 60 hidden layer model')
plt.xlabel('Iteration (in 200s)')
plt.ylabel('Loss')

plt.legend()
plt.savefig('training_vs_validation_loss_rsimple.png')
plt.show()
```



```
[ ]: y_train_pred_simple = []
      for x in X_train:
          y_pred = simple_model.predict(x)
          y_train_pred_simple.append(y_pred)

      y_train_pred_simple = np.array(y_train_pred_simple)
      print(classification_report(y_train, y_train_pred_simple))
```

	precision	recall	f1-score	support
0	0.98	0.66	0.79	4968
1	0.24	0.99	0.39	5689
2	0.90	0.55	0.68	4967
3	0.87	0.36	0.51	5141
4	0.47	0.79	0.59	4872
5	0.03	0.00	0.00	4524
6	1.00	0.12	0.21	4932
7	0.96	0.72	0.83	5192
8	0.01	0.00	0.00	4832
9	0.20	0.18	0.19	4883
accuracy			0.45	50000
macro avg	0.57	0.44	0.42	50000
weighted avg	0.57	0.45	0.43	50000

20 Using packages to mimic the same network

```
[4]: import torch.nn as nn
      import torch.optim as optim
      from torch.utils.data import DataLoader, TensorDataset

      class MLP(nn.Module):
          def __init__(self, input_size, hidden_layers, output_size, g=nn.Sigmoid(),
              ↪0=nn.Softmax(dim=1)):
              super(MLP, self).__init__()
              layers = []
              layer_sizes = [input_size] + hidden_layers
              for i in range(len(layer_sizes) - 1):
                  layers.append(nn.Linear(layer_sizes[i], layer_sizes[i + 1]))
                  layers.append(g)

              layers.append(nn.Linear(layer_sizes[-1], output_size))
              layers.append(0)

              self.model = nn.Sequential(*layers)
```

```
def forward(self, x):
    return self.model(x)
```

21 Trial Run

```
[5]: import torch
from torch.utils.data import DataLoader, random_split

batch_size = 64

train_size = int(0.8 * len(training_data))
val_size = len(training_data) - train_size
train_data, val_data = random_split(training_data, [train_size, val_size])

train_loader = DataLoader(dataset=train_data, batch_size=batch_size,
    ↪shuffle=True)
val_loader = DataLoader(dataset=val_data, batch_size=batch_size, shuffle=False)
test_loader = DataLoader(dataset=test_data, batch_size=batch_size,
    ↪shuffle=False)

[6]: def train(model, train_loader, val_loader, criterion, optimizer, epochs,
    ↪input_size=28*28, device='cpu'):
    """
    Trains the given model on the provided dataset using GPU if available.

    Args:
    - model: The neural network model (an instance of nn.Module).
    - train_loader: DataLoader for the training dataset.
    - val_loader: DataLoader for the validation/test dataset.
    - criterion: Loss function.
    - optimizer: Optimizer.
    - epochs: Number of epochs to train the model.
    - input_size: Size of the flattened input (e.g., 28*28 for MNIST).
    - device: Device to use for computation ('cpu' or 'cuda').

    Returns:
    - train_losses: List of training losses per epoch.
    - val_losses: List of validation losses per epoch.
    """

    model.to(device)
```

```

train_losses = []
val_losses = []

for epoch in range(epochs):
    model.train()
    running_loss = 0.0

    for images, labels in train_loader:

        images, labels = images.view(-1, input_size).to(device), labels.
        ↪to(device)

        outputs = model(images)
        loss = criterion(outputs, labels)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

    train_loss = running_loss / len(train_loader)
    train_losses.append(train_loss)

    model.eval()
    val_loss = 0.0
    with torch.no_grad():
        for images, labels in val_loader:
            images, labels = images.view(-1, input_size).to(device), labels.
            ↪to(device)

            outputs = model(images)
            loss = criterion(outputs, labels)
            val_loss += loss.item()

    val_loss = val_loss / len(val_loader)
    val_losses.append(val_loss)

    print(f"Epoch [{epoch+1}/{epochs}], Training Loss: {train_loss:.4f}, ↪
    ↪Validation Loss: {val_loss:.4f}")

return train_losses, val_losses

```

```
[8]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

baseline_torch_model = MLP(input_size=28*28, hidden_layers=[500, 250, 100],
    ↪output_size=10)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(baseline_torch_model.parameters(), lr=0.01)

baseline_torch_model
```

```
[8]: MLP(
  (model): Sequential(
    (0): Linear(in_features=784, out_features=500, bias=True)
    (1): Sigmoid()
    (2): Linear(in_features=500, out_features=250, bias=True)
    (3): Sigmoid()
    (4): Linear(in_features=250, out_features=100, bias=True)
    (5): Sigmoid()
    (6): Linear(in_features=100, out_features=10, bias=True)
    (7): Softmax(dim=1)
  )
)
```

```
[9]: optimizer.defaults
```

```
[9]: {'lr': 0.01,
      'betas': (0.9, 0.999),
      'eps': 1e-08,
      'weight_decay': 0,
      'amsgrad': False,
      'maximize': False,
      'foreach': None,
      'capturable': False,
      'differentiable': False,
      'fused': None}
```


```
[11]: train_losses, val_losses = train(baseline_torch_model, train_loader,
    ↪val_loader, criterion, optimizer, epochs=15, device=device)
```

```
Epoch [1/15], Training Loss: 1.7528, Validation Loss: 1.5491
Epoch [2/15], Training Loss: 1.5431, Validation Loss: 1.5396
Epoch [3/15], Training Loss: 1.5271, Validation Loss: 1.5226
Epoch [4/15], Training Loss: 1.5223, Validation Loss: 1.5144
Epoch [5/15], Training Loss: 1.5182, Validation Loss: 1.5173
Epoch [6/15], Training Loss: 1.5156, Validation Loss: 1.5175
Epoch [7/15], Training Loss: 1.5158, Validation Loss: 1.5130
Epoch [8/15], Training Loss: 1.5122, Validation Loss: 1.5107
```


Epoch [9/15], Training Loss: 1.5115, Validation Loss: 1.5149
Epoch [10/15], Training Loss: 1.5124, Validation Loss: 1.5137
Epoch [11/15], Training Loss: 1.5104, Validation Loss: 1.5134
Epoch [12/15], Training Loss: 1.5100, Validation Loss: 1.5194
Epoch [13/15], Training Loss: 1.5127, Validation Loss: 1.5137
Epoch [14/15], Training Loss: 1.5076, Validation Loss: 1.5134
Epoch [15/15], Training Loss: 1.5082, Validation Loss: 1.5209

21.1 Plotting a train and validation curve

```
[12]: plt.plot(range(1, len(train_losses)+1), train_losses, label='Training Loss')  
plt.plot(range(1, len(val_losses)+1), val_losses, label='Validation Loss')  
plt.xlabel('Epoch')  
plt.ylabel('Loss')  
plt.title('Training and Validation Loss')  
plt.legend()  
plt.show()
```



```
[25]: def evaluate_data_with_predictions(model, data_loader, criterion,
    ↪ input_size=28*28, device='cpu'):
    """
    Evaluates the model on the entire dataset and returns predictions, test
    ↪ loss, and accuracy.

    Args:
    - model: The trained neural network model.
    - test_loader: DataLoader for the test dataset.
    - criterion: The loss function to calculate the test loss.
    - input_size: Size of the flattened input (e.g., 28*28 for MNIST).
    - device: Device to use for computation ('cpu' or 'cuda').

    Returns:
    - avg_test_loss: The average test loss over the test dataset.
    - accuracy: The accuracy of the model on the test dataset.
    - all_predictions: The list of all predicted labels for the test dataset.
    - all_true_labels: The list of all true labels from the test dataset.
    """
    model.eval()
    model.to(device)

    test_loss = 0.0
    correct = 0
    total = 0
    all_predictions = []
    all_true_labels = []

    with torch.no_grad():
        for images, labels in data_loader:

            images, labels = images.view(-1, input_size).to(device), labels.
            ↪ to(device)

            outputs = model(images)
            loss = criterion(outputs, labels)
            test_loss += loss.item()
            _, predicted = torch.max(outputs, 1)
            all_predictions.extend(predicted.cpu().numpy())
            all_true_labels.extend(labels.cpu().numpy())

            correct += (predicted == labels).sum().item()
            total += labels.size(0)

    return all_predictions, all_true_labels
```

```
[26]: y_pred_train, y_true_train =
      ↪ evaluate_data_with_predictions(baseline_torch_model, train_loader,
      ↪ criterion=criterion, device=device)
      print(classification_report(y_true_train, y_pred_train))
```

	precision	recall	f1-score	support
0	0.98	0.97	0.97	4768
1	0.97	0.99	0.98	5410
2	0.97	0.94	0.95	4793
3	0.99	0.87	0.93	4953
4	0.92	0.96	0.94	4558
5	0.95	0.96	0.95	4323
6	0.92	0.99	0.95	4710
7	0.97	0.95	0.96	5044
8	0.94	0.89	0.91	4755
9	0.86	0.93	0.89	4686
accuracy			0.94	48000
macro avg	0.95	0.94	0.94	48000
weighted avg	0.95	0.94	0.94	48000

```
[27]: y_pred_test, y_true_test = evaluate_data_with_predictions(baseline_torch_model,
      ↪ test_loader, criterion=criterion, device=device)
      print(classification_report(y_true_test, y_pred_test))
```

	precision	recall	f1-score	support
0	0.97	0.98	0.97	980
1	0.97	0.99	0.98	1135
2	0.97	0.94	0.96	1032
3	0.98	0.88	0.92	1010
4	0.92	0.97	0.95	982
5	0.94	0.94	0.94	892
6	0.90	0.97	0.93	958
7	0.96	0.93	0.94	1028
8	0.95	0.88	0.92	974
9	0.86	0.93	0.89	1009
accuracy			0.94	10000
macro avg	0.94	0.94	0.94	10000
weighted avg	0.94	0.94	0.94	10000